

Reading list — SENTINEL

Where to Look: Audit-Allocation Games against Adaptive Persistent Poisoning

FSE 2027 Research Track · FSE-2027-15 · proposal P9 (rank 15)

A defender with a limited verification budget, against an attacker who moves to wherever the budget is not.

Read in this order

If you read nothing else, read these, in this sequence. Each line says what you should be able to answer afterwards.

1. **Tambe** [20] — Stackelberg security games; the formulation that must be exceeded.
2. **Incentive-aware AI safety via Stackelberg games** [12] — the direct precedent for the generic framing.
3. **Korzhyk et al.** [14] — complexity of optimal Stackelberg strategies.
4. **HarnessSafe** [10] — persistence across carriers; the source of the multi-carrier state.
5. **AgentPoison** [5] and **MINJA** [8] — the attack categories reused.
6. **MemSecBench** [16] — the poisoning lifecycle the audit stages map onto.

Closest prior work — the papers this must be distinguished from

These occupy the neighbouring territory. The contribution is only defensible relative to them, and a reviewer will know them.

Incentive-aware AI safety via Stackelberg security games [12]

Occupies the generic “limited audit budget against an adaptive attacker” formulation. The paper must add structure beyond this or it is an application.

Security and game theory [20]

Deployed Stackelberg security games; the canonical treatment.

HarnessSafe [10]

Persistent influence across carriers; the empirical basis for multi-carrier state.

MemSecBench [16]

Poisoning lifecycle and selective repair; maps onto the audit stages.

Adaptive Contracts [1]

Two-stage costly evaluation; the nearest mechanism baseline.

Foundations — the machinery being used

Claimed as building blocks, not contributions. Know them well enough to state precisely what is standard and what is not.

Computing optimal randomized resource allocations [13]

Scalable Stackelberg computation.

Complexity of computing optimal Stackelberg strategies [14]

When the problem is tractable.

Security and game theory [20]

Deployment lessons; why randomization is non-negotiable.

Stochastic network optimization [18]

Sequential decision machinery under budget.

Cluster bootstrap [3]

Instances from one workflow are correlated.

Benjamini–Hochberg [2]
Secondary comparisons.

Benchmarks, baselines and evaluation practice

What the evaluation is built on, compared against, or measured with.

AgentPoison [5]
Retrieval-time poisoning; a reused attack category.

MINJA [8]
Memory injection through ordinary interaction.

MemPoison [15]
Indirect persistence.

Sleeper agents [11]
Delayed activation; the empirical basis for the delay axis.

NeuroTaint [19]
Influence auditing; the detector family.

OpenHands [21]
Workflow harness.

Wider context

Useful for framing and for the related-work section; not on the critical path.

Dependency-guided rollback [7]
Recovery after detection; composes with auditing and changes the optimal policy.

MemTxn [17]
Recovery without the physical write set.

Colosseum [6]
Auditing coordinated agent behaviour.

DRAM [9]
Strategic reporting mechanisms.

Poisoning web-scale training datasets is practical [4]
Poisoning feasibility at scale; framing.

Citation caveat. Entries drawn from 2026-dated preprints inherit their metadata from the source audit document and are flagged % **VERIFY** in **refs.bib**. Verify each against DBLP / arXiv / DOI before citing it in a submission; see **PLAN.md**, *Citation integrity gate*.

References

- [1] Adaptive Contracts Authors. Adaptive contracts for costly evaluation in multi-agent systems. *arXiv preprint arXiv:2603.17212*, 2026.
- [2] Yoav Benjamini and Yosef Hochberg. Controlling the false discovery rate: A practical and powerful approach to multiple testing. *Journal of the Royal Statistical Society Series B*, 57(1):289–300, 1995.
- [3] A. Colin Cameron, Jonah B. Gelbach, and Douglas L. Miller. Bootstrap-based improvements for inference with clustered errors. *The Review of Economics and Statistics*, 90(3):414–427, 2008.
- [4] Nicholas Carlini, Matthew Jagielski, Christopher A. Choquette-Choo, Daniel Paleka, Will Pearce, Hyrum Anderson, Andreas Terzis, Kurt Thomas, and Florian Tramèr. Poisoning web-scale training datasets is practical. In *IEEE Symposium on Security and Privacy (S&P)*, 2024.

- [5] Zhaorun Chen, Zhen Xiang, Chaowei Xiao, Dawn Song, and Bo Li. AgentPoison: Red-teaming LLM agents via poisoning memory or knowledge bases. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2024.
- [6] Colosseum Authors. Colosseum: Auditing collusion among language-model agents. *arXiv preprint arXiv:2602.15198*, 2026.
- [7] Dependency-Guided Rollback Authors. Dependency-guided rollback for compromised agent memory and execution state. *arXiv preprint arXiv:2608.10502*, 2026.
- [8] Shen Dong, Shaochen Xu, Pengfei He, Yige Li, Jiliang Tang, Tianming Liu, Hui Liu, and Zhen Xiang. A practical memory injection attack against LLM agents. *arXiv preprint arXiv:2503.03704*, 2025.
- [9] DRAM Authors. DRAM: Deliberative reporting and mechanisms for strategic agents. *arXiv preprint arXiv:2512.21794*, 2025.
- [10] HarnessSafe Authors. HarnessSafe: Tracing persistent influence across agent carriers. *arXiv preprint arXiv:2608.06984*, 2026.
- [11] Evan Hubinger, Carson Denison, Jesse Mu, Mike Lambert, Meg Tong, Monte MacDiarmid, Tamera Lanham, Daniel M. Ziegler, Tim Maxwell, Newton Cheng, Adam Jermyn, Amanda Asbell, Ansh Radhakrishnan, Cem Anil, David Duvenaud, Deep Ganguli, Fazl Barez, Jack Clark, Kamal Ndousse, Kshitij Sachan, Michael Sellitto, Mrinank Sharma, Nova DasSarma, Roger Grosse, Shauna Kravec, Yuntao Bai, Zachary Witten, Marina Favaro, Jan Brauner, Holden Karnofsky, Paul Christiano, Samuel R. Bowman, Logan Graham, Jared Kaplan, Sören Mindermann, Ryan Greenblatt, Buck Shlegeris, Nicholas Schiefer, and Ethan Perez. Sleeper agents: Training deceptive LLMs that persist through safety training. *arXiv preprint arXiv:2401.05566*, 2024.
- [12] Incentive-Aware AI Safety Authors. Incentive-aware AI safety via Stackelberg security games. *arXiv preprint arXiv:2602.07259*, 2026.
- [13] Christopher Kiekintveld, Manish Jain, Jason Tsai, James Pita, Fernando Ordóñez, and Milind Tambe. Computing optimal randomized resource allocations for massive security games. In *International Conference on Autonomous Agents and Multiagent Systems (AAMAS)*, 2009.
- [14] Dmytro Korzhuk, Vincent Conitzer, and Ronald Parr. Complexity of computing optimal Stackelberg strategies in security resource allocation games. In *AAAI Conference on Artificial Intelligence*, 2010.
- [15] MemPoison Authors. MemPoison: Indirect persistence attacks on agent memory. *arXiv preprint arXiv:2605.29960*, 2026.
- [16] MemSecBench Authors. MemSecBench: Benchmarking the poisoning lifecycle and selective repair of agent memory. *arXiv preprint arXiv:2607.27080*, 2026.
- [17] MemTxn Authors. MemTxn: Source-supported memory updates and recovery without the physical write set. *arXiv preprint arXiv:2607.27834*, 2026.
- [18] Michael J. Neely. *Stochastic Network Optimization with Application to Communication and Queueing Systems*. Morgan & Claypool, 2010.
- [19] NeuroTaint Authors. NeuroTaint: Auditing influence propagation in agent memory. *arXiv preprint arXiv:2604.23374*, 2026.
- [20] Milind Tambe. *Security and Game Theory: Algorithms, Deployed Systems, Lessons Learned*. Cambridge University Press, 2011.
- [21] Xingyao Wang, Boxuan Li, Yufan Song, Frank F. Xu, Xiangru Tang, Mingchen Zhuge, Jiayi Pan, Yueqi Song, Bowen Li, Jaskirat Singh, Hoang H. Tran, Fuqiang Li, Ren Ma, Mingzhang Zheng, Bill Qian, Yanjun Shao, Niklas Muennighoff, Yizhe Zhang, Binyuan Hui, Junyang Lin, Robert Brennan, Hao Peng, Heng Ji, and Graham Neubig. OpenHands: An open platform for AI software developers as generalist agents. In *International Conference on Learning Representations (ICLR)*, 2025.